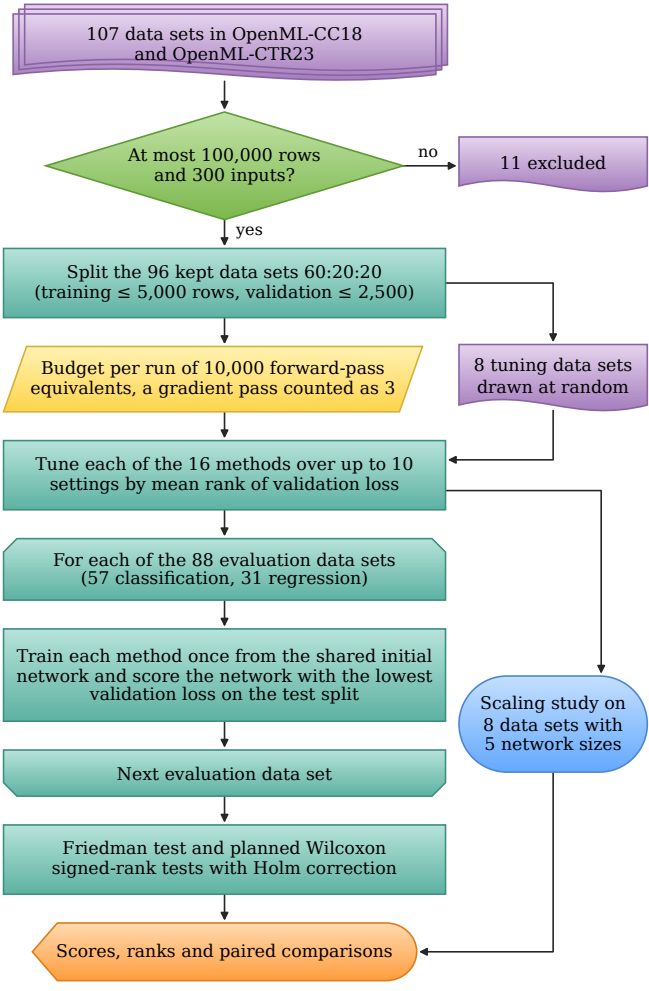
107 data sets in OpenML-CC18 and OpenML-CTR23
At most 100,000 rows and 300 inputs?
no
11 excluded
yes
Split the 96 kept data sets 60:20:20 (training ≤ 5,000 rows, validation ≤ 2,500)
Budget per run of 10,000 forward-pass equivalents, a gradient pass counted as 3
8 tuning data sets drawn at random
Tune each of the 16 methods over up to 10 settings by mean rank of validation loss
For each of the 88 evaluation data sets (57 classification, 31 regression)
Train each method once from the shared initial network and score the network with the lowest validation loss on the test split
Scaling study on 8 data sets with 5 network sizes
Next evaluation data set
Friedman test and planned Wilcoxon signed-rank tests with Holm correction
Scores, ranks and paired comparisons